

Optimizing Data for XGBOOST

Code ▾

This is both a personal guide & my approach to the EY Data Science Challenge 2019:

- Pre-processing data
- Detecting outliers & cleaning data
- Importance of normalisation & balancing in datasets
- Clustering methods
- Tunning using bayesian optimization & xgb.cv
- Training a model
- Evalutation & correlations
- Submissions

Context: Students will have access to a data file data_train.csv that contains the anonymized geolocation data of multiple mobile devices in the City of Atlanta (US) for 11 working days in October 2018. The devices' ID resets every 24 hours; therefore, you will not be able to trace the same device across different days. Therefore, every device ID represents a 1-day journey. Each journey is formed by several trajectories. A trajectory is defined as the route of a moving person in a straight line with an entry and an exit point.

Hide

head(citytrainframe) #I've already changed the id, time, & exit columns on excel.

hash <chr>	id <int>	T... <int>	T... <int>	T... <int>	T... <int>	time_diff <int>	x_entry <dbl>	y_e <
1 0000a8602cf2def930488dee7cdad104_1	0	7	4	7	8	4	3751014	-1909
2 0000a8602cf2def930488dee7cdad104_1	1	7	20	7	25	5	3743937	-1932
3 0000a8602cf2def930488dee7cdad104_1	2	7	53	8	3	10	3744868	-1929
4 0000a8602cf2def930488dee7cdad104_1	3	8	17	8	37	20	3744880	-1929
5 0000a8602cf2def930488dee7cdad104_1	4	14	38	14	38	0	3744909	-1928
6 0000a8602cf2def930488dee7cdad104_1	5	15	2	15	18	16	3744945	-1928

6 rows | 1-10 of 12 columns

PRE-PROCESSING

What hidden data can we extract?

- After quick look of the data set provided on excell, i noticed that each hash “0467278736f0b4abf9ea7fc75ae634ac_29” ends in the date. Hence i was able to create two new additional pieces of data; Date & Day of the week.
- I also noticed that each trajectory “traj_0467278736f0b4abf9ea7fc75ae634ac_29_8” ended in the number of the trajectory and was able to seperate this into a numeric only column.
- I seperated time “8:16:37” into hour, minutes & time_diff
- I decided to remove the velocity columns as they had too many NA values and would be of no help.

- Distance can be calculated from the two (x1,y1) & (x2,y2) coordinates, which also means velocity can be calculated. You will noticed that i removed this data later. I will explain towards the training section.
- Lastly i removed the x_exit & y_exit of every trajectory at hour 15 & 16 so that both the test and train sets are void of bias. BALANCE!

Can you add more data?

- YES! In traffic engineering, roads use something called a k-factor to determine traffic flow. Luckily Georgia Atlanta has this information publically available on their GDOT API, however, our data does not have longitudes & latitudes which means unfortunately i could not sync this data.
- Weather data could have also been added, however, there was no snow in the month of October 2018 and the rain was less than 0.1mm.
- Possibly by turning the data into a widedata set by hash!

[Hide](#)

```
## LIBRARIES
suppressPackageStartupMessages({ library(tidyverse) # Clean & Visual
library(data.table) })# Clean & Read

setwd("C:/Users/ssoma/Desktop/EY Challenge")
city_test <- fread(file = "keras_test3.csv") #Using fread for faster read times as it was a
large dataset
city_train <- fread(file = "keras_train3.csv")

# Creating Dataframe
citytrainframe <- data.frame(city_train, stringsAsFactors = TRUE)
citytestframe <- data.frame(city_test, stringsAsFactors = TRUE)
citytrainframe <- citytrainframe[,-c(1)]
citytestframe <- citytestframe[,-c(1)]

# Combining Train and Test Data to perform analysis on.
kclusterdata <- rbind(citytrainframe,citytestframe)

# Seperating The date from Hash
kclusterdata <- kclusterdata %>% separate(hash, into = c("hash", "date"), sep = "_")
kclusterdata$date <- as.numeric(kclusterdata$date)
kclusterdata$hash <- as.factor(kclusterdata$hash)
kclusterdata$hash <- as.numeric(kclusterdata$hash)

# Creating a Day column from the Date Column
kclusterdata$day <- cut(kclusterdata$date, breaks = c(-Inf, 1,3,5,9,11,15,19,23,25,29,31), la
bels = c(1,3,5,2,4,1,5,2,4,1,3))
kclusterdata <- kclusterdata[,c(1,14,2:13)]
kclusterdata$day <- as.numeric(kclusterdata$day)

summary(kclusterdata)
```

hash	day	date	id	TENH	TENM	TEXH	TEXM	time_diff	
Min. : 1	Min. :1.000	Min. : 1.00	Min. : 0.000	Min. : 0.00	Min. :	0.00	Min. : 0.00	Min. : 0.00	Min. : 0.000
1st Qu.: 41966	1st Qu.:1.000	1st Qu.: 5.00	1st Qu.: 2.000	1st Qu.: 8.00	1st Qu.:1	4.00	1st Qu.: 8.00	1st Qu.:14.00	1st Qu.: 0.000
Median : 84010	Median :3.000	Median :15.00	Median : 4.000	Median :11.00	Median :2	9.00	Median :11.00	Median :29.00	Median : 0.000
Mean : 83909	Mean :2.846	Mean :16.34	Mean : 5.706	Mean :10.43	Mean :2	9.24	Mean :10.53	Mean :29.28	Mean : 5.779
3rd Qu.:125844	3rd Qu.:4.000	3rd Qu.:25.00	3rd Qu.: 8.000	3rd Qu.:14.00	3rd Qu.:4	4.00	3rd Qu.:14.00	3rd Qu.:44.00	3rd Qu.: 8.000
Max. :167578	Max. :5.000	Max. :31.00	Max. :66.000	Max. :16.00	Max. :5	9.00	Max. :16.00	Max. :59.00	Max. :409.000
x_entry	y_entry	x_exit	y_exit	target					
Min. :3741027	Min. :-19382915	Min. :3740998	Min. :-19376883	Min. :0.000					
1st Qu.:3755179	1st Qu.: -19274682	1st Qu.:3755443	1st Qu.: -19271796	1st Qu.:0.000					
Median :3760079	Median : -19230184	Median :3760075	Median : -19230377	Median :0.000					
Mean :3760406	Mean : -19221509	Mean :3760416	Mean : -19221874	Mean :0.299					
3rd Qu.:3767500	3rd Qu.: -19169903	3rd Qu.:3767361	3rd Qu.: -19172540	3rd Qu.:1.000					
Max. :3777099	Max. : -19042657	Max. :3777055	Max. : -19045740	Max. :1.000					
		NA's :167578	NA's :167578	NA's :16461					

Hide

kclusterdata

hash	day	date	id	TENH	TENM	TEXH	TEXM						time_diff						x_entry					
<dbl>	<dbl>	<dbl>	<int>		<int>		<int>	<int>						<int>						<dbl>				
1	1	1	0		7		4	7	8						4						3751014			
1	1	1	1		7		20	7	25						5						3743937			
1	1	1	2		7		53	8	3						10						3744868			
1	1	1	3		8		17	8	37						20						3744880			
1	1	1	4		14		38	14	38						0						3744909			
1	1	1	5		15		2	15	18						16						3744945			
2	4	9	0		14		29	14	29						0						3749450			
2	4	9	1		14		39	14	39						0						3749090			
2	4	9	2		14		50	14	50						0						3749042			
2	4	9	3		15		0	15	29						29						3749088			
1-10 of 1,017,199 rows 1-10 of 14 columns														Previous		1	2	3	4	5	6	...	100	Next

Hide

NA

OUTLIERS

Univariate Outlier Detection Methods:

- **Tukey’s method of outlier detection:** Tukey’s rule says that the outliers are values more than 1.5 times the interquartile range from the quartiles — either below $Q1 - 1.5IQR$, or above $Q3 + 1.5IQR$.
- **z-score:** Using the outliers package. “z” calculates normal scores (differences between each value and the mean divided by sd).

Multivariate Outlier Detection Methods:

- **Bivariate box plots and scatter plots:** same as above but with more variables.
- **The Mahalanobis distance:** Using the MVN package.

Using trail and error i applied all the above to various aspects of the data to determine outliers. I found Mahalanobis distance to give the best results.

Hide

```
#library(MVN)
# Visualisation
## mvn(data = kclusterdata, multivariateOutlierMethod = "quan", showOutliers = TRUE)
## mvn(data = kclusterdata, multivariateOutlierMethod = "quan", multivariatePlot = "contour")

# Finding Mahalanobis Distance with the data with no NA values!
MAH1 <- kclusterdata[,c(1:6)]
MAH2 <- mahalanobis(MAH1, colMeans(MAH1), cov(MAH1))
kclusterdata$MAH <- round(MAH2, 5)

# Assigning 1 and 0 to data which i think is valid
kclusterdata$outlier_maha <- 0
kclusterdata$outlier_maha[kclusterdata$MAH > 10.5 ] <- 1 # 10.5 was decided via trial and err
or

#no_outliers <- kclusterdata[!(kclusterdata$outlier_maha== 1),]

head(kclusterdata)
```

	hash	day	date	id	TENH	TENM	TEXH	TEXM	time_diff
	<dbl>	<dbl>	<dbl>	<int>	<int>	<int>	<int>	<int>	<int>
1	1	1	1	0	7	4	7	8	4
2	1	1	1	1	7	20	7	25	5
3	1	1	1	2	7	53	8	3	10
4	1	1	1	3	8	17	8	37	20
5	1	1	1	4	14	38	14	38	0
6	1	1	1	5	15	2	15	18	16

6 rows | 1-10 of 16 columns

Hide

NA

Handling Outliers:

Excluding: `df[is.na(df),] <- NULL`

Imputing: Look up guides on mice, missForest, impute, missRanger. This approach may not always be best, sometimes its best to leave values as Na's. Do a test of imputed data vs non-imputed to see what works for you.

Capping: NA

In this project because i assigned 1 and 0s to the Maha values therefore i had no need to remove any data. I did however, attempt to use missRanger and Mice to impute the values in a seperate trial where i combined the orginial train and test data and imputed the NA values in test but, the results were poor.

Feature extraction:

Principal Component Analysis (PCA): prcomp, caret, mlr are useful packages for this. `preProcess(df, method = "pca", thresh = 0.90)` – cumulative explained variance of 90%. `preProcess(df, method = "pca", pcaComp = number_of_dimensions_if_known)`

I did attempt some PCA, however, this showed no improvement in results and hence didn't end up using it.

NORMALISATION

Pick the right nromalisation method for your data.

Centering and scaling: use the `scale()` function. For mean centering use (`center = TRUE, scale = FALSE`). For scaling (`center = FALSE, scale = TRUE`). For RMS (`center = FALSE, scale = apply(df, 2, sd, na.rm = TRUE)`)

z score standardisation: using the `scale()` with the `center = TRUE, scale = TRUE`

Min- Max Normalisation: Feel free to use this in combination with `lapply()` or pipes. `minmaxnormalise <- function(x) {(x-min(x,na.rm = TRUE))/(max(x, na.rm =TRUE)-min(x, na.rm = TRUE))}`

Data transformation for skewed data.

Box-Cox Transformation: The caret, MASS, forecast, geoR, EnvStats, and AIS packages can all perform BCT. I personally like Forecast as it has functions to find the best lambda parameter.

Data Transformation via Mathematical Operations: Use `log10()`, `log()`

Binning/ Discretisation

Equal width (distance) binning: Use the `infotheo` package and the `discretize()` function

Equal depth (frequency) binning: `discretize()` function with `disc = "equalfreq"`

I used all the above methods when it came to normalising the data thorough trial and error. I found that min-max, centering and scaling were the best options.

Hide

```
options(na.action='na.pass')
minmaxnormalise <- function(x) {(x-min(x,na.rm = TRUE))/(max(x, na.rm =TRUE)-min(x, na.rm = T
RUE))}

kclusterdata$xEnNorm <- kclusterdata$x_entry %>% minmaxnormalise()
kclusterdata$yEnNorm <- kclusterdata$y_entry %>% minmaxnormalise()
kclusterdata$timeNorm <- kclusterdata$time_diff %>% minmaxnormalise()

head(kclusterdata)
```

	hash <dbl>	day <dbl>	date <dbl>	id <int>	TENH <int>	TENM <int>	TEXH <int>	TEXM <int>	time_diff <int>	
1	1	1	1	0	7	4	7	8	4	
2	1	1	1	1	7	20	7	25	5	
3	1	1	1	2	7	53	8	3	10	
4	1	1	1	3	8	17	8	37	20	
5	1	1	1	4	14	38	14	38	0	
6	1	1	1	5	15	2	15	18	16	

6 rows | 1-10 of 19 columns

Hide

NA

CLUSTERING

Pick the right clustering method for your data.

K-means Clustering:

Hierarchical clustering:

Expectation Maximization Clustering:

Fuzzy clustering:

Density-based clustering:

Model-based clustering:

Partitioning methods:

Source & read more (<https://www.datanovia.com/en/blog/types-of-clustering-methods-overview-and-quick-start-r-code/>)

A detailed explanation of these clustering methods are in the link above. This section of the project was by far my favourite. Clustering methods are so facinating especially when you can visualise how they work.

Of all the clustering methods k-means and EM Clustering seem to do the best job. The visual printed below doesnt show much detail, but in the high quality saved file you can see that that the clusters have a bias towards the x-axis. Would love to know more about how others have applied clustering.

theme_black() credit (<https://jonlefccheck.net/2013/03/11/black-theme-for-ggplot2-2/>)

[Hide](#)

```
# Applying Clustering to the entry points
clustertrain <- kmeans(kclusterdata[,10:11], 5000, iter.max=40)
kclusterdata$cluster <- as.numeric(clustertrain$cluster)
summary(kclusterdata)
```

hash	day	date	id	TENH	TENM
TEXH	TEXM	time_diff			
Min. : 1	Min. :1.000	Min. : 1.00	Min. : 0.000	Min. : 0.00	Min. :
0.00	Min. : 0.00	Min. : 0.00	Min. : 0.000		
1st Qu.: 41966	1st Qu.:1.000	1st Qu.: 5.00	1st Qu.: 2.000	1st Qu.: 8.00	1st Qu.:1
4.00	1st Qu.: 8.00	1st Qu.:14.00	1st Qu.: 0.000		
Median : 84010	Median :3.000	Median :15.00	Median : 4.000	Median :11.00	Median :2
9.00	Median :11.00	Median :29.00	Median : 0.000		
Mean : 83909	Mean :2.846	Mean :16.34	Mean : 5.706	Mean :10.43	Mean :2
9.24	Mean :10.53	Mean :29.28	Mean : 5.779		
3rd Qu.:125844	3rd Qu.:4.000	3rd Qu.:25.00	3rd Qu.: 8.000	3rd Qu.:14.00	3rd Qu.:4
4.00	3rd Qu.:14.00	3rd Qu.:44.00	3rd Qu.: 8.000		
Max. :167578	Max. :5.000	Max. :31.00	Max. :66.000	Max. :16.00	Max. :5
9.00	Max. :16.00	Max. :59.00	Max. :409.000		

x_entry	y_entry	x_exit	y_exit	target
MAH	outlier_maha	xEnNorm	yEnNorm	
Min. :3741027	Min. :-19382915	Min. :3740998	Min. :-19376883	Min. :0.000
Min. : 0.1065	Min. :0.00000	Min. :0.0000	Min. :0.0000	
1st Qu.:3755179	1st Qu.: -19274682	1st Qu.:3755443	1st Qu.: -19271796	1st Qu.:0.000
1st Qu.: 4.0991	1st Qu.:0.00000	1st Qu.:0.3923	1st Qu.:0.3181	
Median :3760079	Median : -19230184	Median :3760075	Median : -19230377	Median :0.000
Median : 5.4787	Median :0.00000	Median :0.5281	Median :0.4489	
Mean :3760406	Mean : -19221509	Mean :3760416	Mean : -19221874	Mean :0.299
Mean : 6.0000	Mean :0.06221	Mean :0.5372	Mean :0.4744	
3rd Qu.:3767500	3rd Qu.: -19169903	3rd Qu.:3767361	3rd Qu.: -19172540	3rd Qu.:1.000
3rd Qu.: 7.1222	3rd Qu.:0.00000	3rd Qu.:0.7339	3rd Qu.:0.6260	
Max. :3777099	Max. : -19042657	Max. :3777055	Max. : -19045740	Max. :1.000
Max. :133.3027	Max. :1.00000	Max. :1.0000	Max. :1.0000	
		NA's :167578	NA's :167578	NA's :16461

timeNorm	cluster
Min. :0.00000	Min. : 1
1st Qu.:0.00000	1st Qu.:1250
Median :0.00000	Median :2510
Mean :0.01413	Mean :2504
3rd Qu.:0.01956	3rd Qu.:3763
Max. :1.00000	Max. :5000

[Hide](#)

```

library(gridExtra)
# Black theme
theme_black = function(base_size = 12, base_family = "") {

  theme_grey(base_size = base_size, base_family = base_family) %+replace%

  theme(
    # Specify axis options
    axis.line = element_blank(),
    axis.text.x = element_text(size = base_size*0.8, color = "white", lineheight = 0.9),
    axis.text.y = element_text(size = base_size*0.8, color = "white", lineheight = 0.9),
    axis.ticks = element_line(color = "white", size = 0.2),
    axis.title.x = element_text(size = base_size, color = "white", margin = margin(0, 10, 0
, 0)),
    axis.title.y = element_text(size = base_size, color = "white", angle = 90, margin = mar
gin(0, 10, 0, 0)),
    axis.ticks.length = unit(0.3, "lines"),
    # Specify legend options
    legend.background = element_rect(color = NA, fill = "black"),
    legend.key = element_rect(color = "white", fill = "black"),
    legend.key.size = unit(1.2, "lines"),
    legend.key.height = NULL,
    legend.key.width = NULL,
    legend.text = element_text(size = base_size*0.8, color = "white"),
    legend.title = element_text(size = base_size*0.8, face = "bold", hjust = 0, color = "wh
ite"),
    legend.position = "right",
    legend.text.align = NULL,
    legend.title.align = NULL,
    legend.direction = "vertical",
    legend.box = NULL,
    # Specify panel options
    panel.background = element_rect(fill = "black", color = NA),
    panel.border = element_rect(fill = NA, color = "white"),
    panel.grid.major = element_line(color = "grey35"),
    panel.grid.minor = element_line(color = "grey20"),
    panel.spacing = unit(0.5, "lines"),
    # Specify facetting options
    strip.background = element_rect(fill = "grey30", color = "grey10"),
    strip.text.x = element_text(size = base_size*0.8, color = "white"),
    strip.text.y = element_text(size = base_size*0.8, color = "white", angle = -90),
    # Specify plot options
    plot.background = element_rect(color = "black", fill = "black"),
    plot.title = element_text(size = base_size*1.2, color = "white"),
    plot.margin = unit(rep(1, 4), "lines")

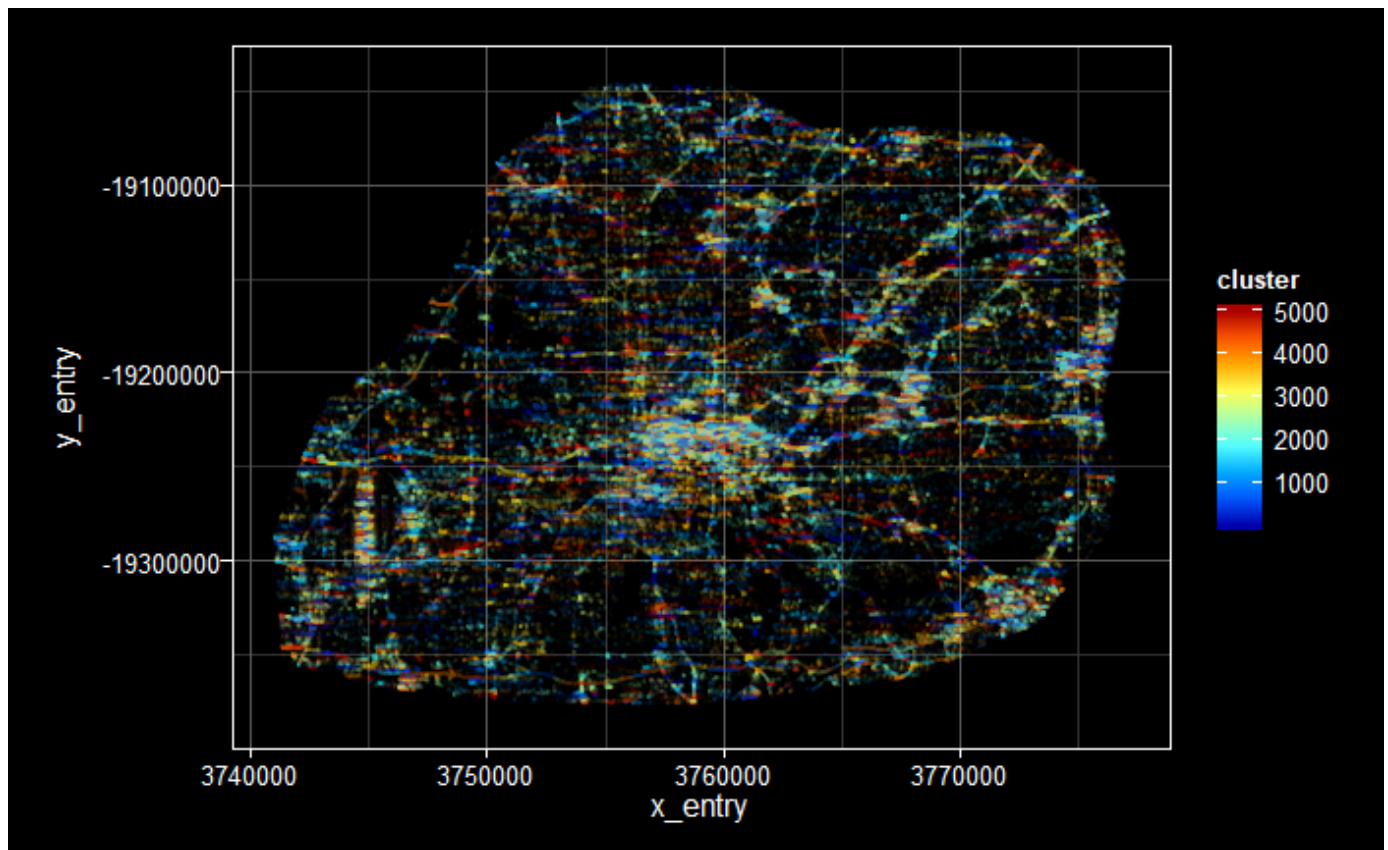
  )

}
library(colorRamps)
#Visualisation of the Clusters
ggplot(kclusterdata) + theme_black() +
  aes(x_entry, y_entry, color = cluster) +
  geom_point(alpha = 0.008, size = 0.0001) +
  scale_color_gradientn(colours=matlab.like(5000))
  #scale_colour_viridis(option = "C") # from the viridis package

```



```
ggsave("kcluster5000black.jpg", units="in", width=14, height=10, dpi=1500)
```

[Hide](#)

NA
NA

TUNING METHOD: BAYESIAN OPTIMIZATION

There are many tuning methods i've tried for this project. Bayesian Optimization with mlrMBO turned out to be the most efficient and accurate.

mlrMBO

[5.1] Bischl B, Richter J, Bossek J, Horn D, Thomas J, Lang M (2017). `_mlrMBO`: A Modular Framework for Model-Based Optimization of Expensive Black-Box Functions_. (<http://arxiv.org/abs/1703.03373>)

[5.2] Code sampled from Simon Coulombe (<https://www.simoncoulombe.com/2019/01/bayesian/>)

[Hide](#)

```
# print a summary with run  
run
```

Recommended parameters:
eta=0.00102; gamma=2.38; max_depth=1; min_child_weight=2; subsample=0.404; colsample_bytree=0.359
Objective: y = 0.707

Optimization path
11 + 10 entries in total, displaying last 10 (or less):

	eta	gamma	max_de...	min_child_weight	subsam...	colsample_bytree	
	<dbl>	<dbl>	<int>	<int>	<dbl>	<dbl>	
12	0.001664730	3.412954	5	4	0.3871472	0.5752815	0.700
13	0.001771330	3.397143	2	8	0.8515646	0.5873521	0.700
14	0.002565296	2.954818	6	10	0.2317204	0.2867190	0.700
15	0.005011954	3.287664	7	8	0.8498699	0.4858637	0.700
16	0.001008299	4.071173	1	5	0.2019193	0.9493196	0.700
17	0.001024414	2.377319	1	2	0.4039261	0.3591744	0.700
18	0.001707212	4.983310	1	1	0.2024028	0.3363391	0.700
19	0.001059178	4.119342	1	2	0.5103523	0.7686407	0.700
20	0.001012248	4.139225	1	7	0.2023848	0.6222008	0.700
21	0.001099040	4.057689	1	4	0.3737862	0.3326670	0.700

1-10 of 10 rows | 1-10 of 19 columns

Hide

```
# return best model hyperparameters using  
run$x
```

```
$eta  
[1] 0.001024414  
  
$gamma  
[1] 2.377319  
  
$max_depth  
[1] 1  
  
$min_child_weight  
[1] 2  
  
$subsample  
[1] 0.4039261  
  
$colsample_bytree  
[1] 0.3591744
```

Hide

```
# return best log likelihood using
run$y
```

[1] 0.7073641

Hide

```
# return all results using
run$opt.path$env$path
```

eta	gamma	max_de...	min_child_weight	subsam...	colsample_bytree			
<dbl>	<dbl>	<int>	<int>	<dbl>	<dbl>			
0.010000000	0.00000000	6	3	0.8000000	0.8000000	0.70		
0.029653675	2.05908861	6	6	0.6746934	0.7706236	0.69		
0.023476226	4.40334366	10	8	0.5701985	0.4679024	0.69		
0.042117124	4.92683403	4	1	0.8041056	0.3855289	0.69		
0.033275426	0.07700364	5	4	0.7322089	0.6985217	0.69		
0.001317531	3.05007313	9	10	0.8501541	0.9457132	0.70		
0.048554996	1.00709307	2	5	0.3517939	0.5348991	0.69		
0.037636391	1.86700871	8	9	0.3628718	0.2511837	0.69		
0.007812635	0.54903425	7	2	0.9457862	0.3572049	0.70		
0.013350092	3.82396167	1	3	0.2494433	0.6396891	0.70		
1-10 of 21 rows				Previous	1	2	3	Next

TRAINING A MODEL

Using the best parameters from above.

Hide

```
start.time <- Sys.time()
bst_model <- xgb.train(
  data = dtrain,
  nrounds = 100,
  objective = "binary:logistic",
  booster = "gbtree",
  eval_metric = "error",
  # watchlist = watchlist,
  nfold = 10,
  eta = 0.001099040 ,
  max.depth = 1 ,
  min_child_weight= 4,
  gamma = 4.057689, #makes it more conservation, to avoid overfitting
  subsample = 0.3737862, #lower values prevents overfitting
  colsample_bytree = 0.3326670 ,
  missing = NA,
  seed = 333)
end.time <- Sys.time()
time.taken <- end.time - start.time
time.taken
```

Time difference of 17.62415 secs

EVALUATION

Hide

```
# Training & test error plot
e <- data.frame(bst_model$evaluation_log)
#plot(e$iter, e$train_mlogloss, col = 'blue')
#lines(e$iter, e$test_mlogloss, col = 'red')

# Feature importance
imp <- xgb.importance(colnames(dtrain), model = bst_model)
print(imp)
```

Feature <chr>	Gain <dbl>	Cover <dbl>	Frequency <dbl>
y_exit	4.370654e-01	0.3189249290	0.3198
yEnNorm	1.808232e-01	0.1626566431	0.1509
y_entry	1.725442e-01	0.1535352901	0.1421
x_exit	1.202370e-01	0.1588621476	0.1735
x_entry	4.905289e-02	0.0817891288	0.0858
xEnNorm	3.757046e-02	0.0624352387	0.0653
TEXH	1.563266e-03	0.0251088152	0.0266
TENH	6.677610e-04	0.0154207922	0.0159
id	2.104585e-04	0.0082479186	0.0061


Feature <chr>	Gain <dbl>	Cover <dbl>	Frequency <dbl>
time_diff	1.412149e-04	0.0053504343	0.0066

1-10 of 14 rows

Previous12Next

Hide

```
xgb.plot.importance (importance_matrix = imp[1:8]) #top 8
```



SUBMISSION

Hide

```
head(citytestframe)
```

	hash <dbl>	day <dbl>	date <dbl>	id <int>	TENH <int>	TENM <int>	TEXH <int>	TEXM <int>	time_diff <int>
814263	8	2	31	0	11	43	11	50	7
814264	8	2	31	2	12	21	12	21	0
814265	8	2	31	3	12	34	13	14	40
814266	8	2	31	4	13	25	13	43	18
814267	8	2	31	5	15	3	15	10	7
814268	11	5	25	0	8	8	8	20	12

6 rows | 1-10 of 20 columns

```
btestm <- sparse.model.matrix(target~.-1,citytestframe)
btest_label <- citytestframe[, "target"]
btest_matrix <- xgb.DMatrix(data = (btestm), label = btest_label)
target <- predict(bst_model, newdata = btest_matrix)
sub1 <- read.csv("data_test.csv")
sub2 <- cbind(sub1[,c(3,11)],target)
sub3 <- sub2[(is.na(sub2$x_exit)),]
write.csv(sub3[,c(1,3)], file = "xgbpredNEW7.csv")
sub3[,c(1,3)]
```

	trajectory_id<fctr>	target<dbl>
5	traj_00032f51796fd5437b238e3a9823d13d_31_5	0.11506353
8	traj_000479418b5561ab694a2870cc04fd43_25_10	0.21788977
11	traj_000506a39775e5bca661ac80e3f466eb_29_5	0.62713480
14	traj_0005401ceddaf27a9b7f0d42ef1fbe95_1_4	0.25364435
18	traj_00063a4f6c12e1e4de7d876580620667_3_4	0.19175792
24	traj_0006535be25bb52dd06983447880c964_5_12	0.11517836
28	traj_0006f84bb33ec929d1cda7686f861d0a_31_3	0.62509215
32	traj_00093ae562586aed0e053b8431e8ace4_23_10	0.18646917
35	traj_000c739e444a70e1804d757a0580caaa_31_3	0.62505734
40	traj_000d479078af08618bddc7f09082b8c3_11_6	0.26399857
1-10 of 33,515 rows		Previous 1 2 3 4 5 6 ... 100 Next

Hide

NA

APPENDIX

TUNING METHOD: NESTED LOOP

I used this method over night just to make sure that Bayesian didn't already give me the best results. This method is obviously very time consuming but useful if you know what range you're looking for.

Hide

```

data <- (citytrainframe)
train <- data[1:20]
options(na.action='na.pass')

# binarize all factors
library(caret)
library(Metrics)
dmy <- dummyVars(" ~ .", data = train)
pTrsf <- data.frame(predict(dmy, newdata = train))
#####

# what we're trying to predict adults that make more than 50k
outcomeName <- c('target')
# list of features
predictors <- names(pTrsf)[!names(pTrsf) %in% outcomeName]

# take first 10% of the data only! Depending on your computing power
trainPortion <- floor(nrow(pTrsf)*0.1)

trainSet <- pTrsf[ 1:floor(trainPortion/2),]
testSet <- pTrsf[(floor(trainPortion/2)+1):trainPortion,]

#add eta, gamma & subsample
start.time <- Sys.time()
smallestError <- 100
for (fold in seq(6,32,1)) { #7,15,1
  for (round in seq(1,1000,1)) { # 20,50,1
    for (etax in seq(0.001,0.25,0.001)) { #0.00,0.3, 0.02
      # train
      bst <- xgboost(data = as.matrix(trainSet[,predictors]),
                     label = trainSet[,outcomeName],
                     max.depth=8, nround=4,
                     objective = "binary:logistic", booster = "gbtree",
                     eval_metric = "error", eta = 0.038,
                     missing = NA, # max.depth = 10,
                     min_child_weight= 8, nfold = fold,
                     gamma = 0.0472, #makes it more conservation, to avoid overfitt
ing
                     subsample = 0.698, #lower values prevents overfitting
                     colsample_bytree = 0.6298,
                     seed = 333, verbose=0)

      gc()

      # predict
      predictions <- predict(bst, as.matrix(testSet[,predictors]), outputmargin=TRUE)

      err <- rmse(as.numeric(testSet[,outcomeName]), as.numeric(predictions))

      if (err < smallestError) {
        smallestError = err
        print(paste(fold,err))
      }
    }
  }
}
}

```

```
end.time <- Sys.time()
time.taken <- end.time - start.time
time.taken
```

WIDE DATA TRAINING

This method of turning the dataset into a wide dataset by hash proved to show no significant improvement. However, i believe the reason for this is due to the amount of NA values in the data set.

hash <dbl>	day <dbl>	date <dbl>	A <int>	B <int>	C <int>	D <int>	E <int>	F <int>	G <int>	
1	1	1	1	7	7	7	8	14	15	
1	1	1	2	7	7	7	8	14	15	
1	1	1	3	7	7	7	8	14	15	
1	1	1	1	7	7	7	8	14	15	
1	1	1	1	7	7	7	8	14	15	
1	1	1	1	7	7	7	8	14	15	
2	4	9	1	14	14	14	15	NA	NA	
2	4	9	2	14	14	14	15	NA	NA	
2	4	9	3	14	14	14	15	NA	NA	
2	4	9	1	14	14	14	15	NA	NA	
1-10 of 20 rows 1-10 of 292 columns							Previous	1	2	Next